

Aayush Shah

19BCE245

27 October 2022

Big Data Analytics

Practical 4

Aim

Design MapReduce algorithms to take a very large file of integers and produce as output:

- a) The largest integer
- b) The average of all the integers.

AVERAGE

• Mapper

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;
public class WCMapper extends Mapper<LongWritable, Text, IntWritable,
IntWritable> {
    @Override
    public void map(LongWritable key, Text value, Context ctx) throws
IOException {
        String data[] = value.toString().split(" ");
        for (String num : data) {
            int number = Integer.parseInt(num);
            ctx.write(new IntWritable(1), new IntWritable(number));
        }
    }
}
```

• Reducer

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.FloatWritable;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
public class WCReducer extends Reducer<IntWritable, IntWritable, Text,
FloatWritable> {
    @Override
    public void reduce(IntWritable key, Iterator<IntWritable> value,
Context ctx) throws IOException {
        int sum = 0;
        int count=0;
        while (value.hasNext()) {
            IntWritable i = value.next();
            sum += i.get();
            count++;
        }
        float sum1= (float)(sum)/count;
        ctx.write(new Text("AVG"), new FloatWritable(sum1));
    }
}
```

• Driver

```
import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
public class WCDriver {
    public static void main(String args[]) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "max int");
        conf.setMapperClass(WCMapper.class);
        conf.setReducerClass(WCReducer.class);
        conf.setMapOutputKeyClass(IntWritable.class);
        conf.setMapOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

LARGEST

• Mapper

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;
public class WCMapper extends Mapper<LongWritable, Text, IntWritable,
IntWritable> {
    @Override
    public void map(LongWritable key, Text value, Context ctx) throws
IOException {
        String[] data = value.toString().split(" ");
        byte b;
        int i;
        String[] arrayOfString1;
        for (i = (arrayOfString1 = data).length, b = 0; b < i; ) {
            String num = arrayOfString1[b];
            int number = Integer.parseInt(num);
            ctx.write(new IntWritable(number), new IntWritable(1));
            b++;
        }
    }
}
```

• Reducer

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
public class WCReducer extends Reducer<IntWritable, IntWritable,
IntWritable, IntWritable> {
    int max1 = -1000000000;
    @Override
    public void reduce(IntWritable key, Iterator<IntWritable> value, Context
ctx) throws IOException {
        this.max1 = Math.max(this.max1, key.get());
        ctx.write(key, new IntWritable(this.max1));
    }
}
```

• Driver

```
import java.io.IOException;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
```

```

import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
public class WCDriver {
    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "average");
        conf.setMapperClass(WCMapper.class);
        conf.setReducerClass(WCReducer.class);
        conf.setMapOutputKeyClass(IntWritable.class);
        conf.setMapOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

DISTINCT

• Mapper

```

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;
public class WCMapper extends Mapper<LongWritable, Text, IntWritable,
IntWritable> {
    @Override
    public void map(LongWritable key, Text value, Context ctx) throws
IOException {
        String data[] = value.toString().split(" ");
        for (String num : data) {
            int number = Integer.parseInt(num);
            ctx.write(new IntWritable(number), new IntWritable(number));
        }
    }
}

```

• Reducer

```

import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.FloatWritable;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
public class WCReducer extends Reducer<IntWritable, IntWritable,
IntWritable, IntWritable> {
    @Override
    public void reduce(IntWritable key, Iterator<IntWritable> value,
Context ctx) throws IOException {

```

```

    int counter=0;
    while (value.hasNext()) {
        IntWritable i = value.next();
        counter++;
        if(counter==1)
            ctx.write(new IntWritable(i.get()), new
IntWritable(i.get()));
    }
} }

```

• Driver

```

import java.io.IOException;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
public class WCDriver {
    public static void main(String args[]) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "distinct nums");
        conf.setMapperClass(WCMapper.class);
        conf.setReducerClass(WCReducer.class);
        conf.setMapOutputKeyClass(IntWritable.class);
        conf.setMapOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    } }

```

Conclusion

From this practical, I performed map reduce operations for finding largest, average and distinct element in Hadoop.